

SURUTHI S

225229141

```
In [2]: 1 import pandas as pd
        2 import numpy as np
        3 import warnings
        4 warnings.filterwarnings('ignore')
```

```
In [3]: 1 customer_data = pd.read_csv('Mall_Customers.csv')
        2 customer_data.head()
```

```
Out[3]:
```

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [4]: 1 customer_data.shape
```

```
Out[4]: (200, 5)
```

```
In [5]: 1 customer_data.columns
```

```
Out[5]: Index(['CustomerID', 'Genre', 'Age', 'Annual Income (k$)',
              'Spending Score (1-100)'],
              dtype='object')
```

```
In [6]: 1 customer_data.dtypes
```

```
Out[6]: CustomerID          int64
Genre          object
Age           int64
Annual Income (k$)  int64
Spending Score (1-100)  int64
dtype: object
```

```
In [7]: 1 customer_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   CustomerID                  200 non-null   int64
1   Genre                       200 non-null   object
2   Age                         200 non-null   int64
3   Annual Income (k$)          200 non-null   int64
4   Spending Score (1-100)      200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
In [8]: 1 customer_data.value_counts()
```

```
Out[8]: CustomerID  Genre  Age  Annual Income (k$)  Spending Score (1-100)
1             Male    19    15                   39                1
138           Male    32    73                   73                1
128           Male    40    71                   95                1
129           Male    59    71                   11                1
130           Male    38    71                   75                1
..
70           Female   32    48                   47                1
71           Male    70    49                   55                1
72           Female   47    49                   42                1
73           Female   60    50                   49                1
200          Male    30   137                   83                1
Length: 200, dtype: int64
```

Step2 [Label encode gender]

```
In [9]: 1 from sklearn import preprocessing
2 label_encoder = preprocessing.LabelEncoder()
3 customer_data['Genre'] = label_encoder.fit_transform(customer_data['Genre'])
4 customer_data.head()
```

```
Out[9]:   CustomerID  Genre  Age  Annual Income (k$)  Spending Score (1-100)
0           1      1    19                   15                39
1           2      1    21                   15                81
2           3      0    20                   16                 6
3           4      0    23                   16                77
4           5      0    31                   17                40
```

Step3 [Check for variance]

```
In [10]: 1 customer_data.describe()
```

```
Out[10]:   CustomerID      Genre      Age  Annual Income (k$)  Spending Score (1-100)
count  200.000000  200.000000  200.000000      200.000000      200.000000
mean    100.500000    0.440000   38.850000     60.560000     50.200000
std     57.879185    0.497633   13.969007     26.264721     25.823522
min      1.000000    0.000000   18.000000     15.000000      1.000000
25%     50.750000    0.000000   28.750000     41.500000     34.750000
50%    100.500000    0.000000   36.000000     61.500000     50.000000
75%    150.250000    1.000000   49.000000     78.000000     73.000000
max    200.000000    1.000000   70.000000    137.000000     99.000000
```

```
In [11]: 1 customer_data.var()
```

```
Out[11]: CustomerID      3350.000000
Genre          0.247638
Age           195.133166
Annual Income (k$)  689.835578
Spending Score (1-100)  666.854271
dtype: float64
```

In [12]:

1customer_data.corr()

Out[12]:

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
CustomerID	1.000000	0.057400	-0.026763	0.977548	0.013835
Genre	0.057400	1.000000	0.060867	0.056410	-0.058109
Age	-0.026763	0.060867	1.000000	-0.012398	-0.327227
Annual Income (k\$)	0.977548	0.056410	-0.012398	1.000000	0.009903
Spending Score (1-100)	0.013835	-0.058109	-0.327227	0.009903	1.000000

Step4 [Check skewness]

In [13]:

1customer_data.skew()

Out[13]:

CustomerID0.000000
Genre0.243578
Age0.485569
Annual Income (k\$)0.321843
Spending Score (1-100)-0.047220
dtype: float64

In [14]:

1customer_data.sort_values(by=['Genre','Age','Annual Income (k\$)','Spending Score (1-100)'])

Out[14]:

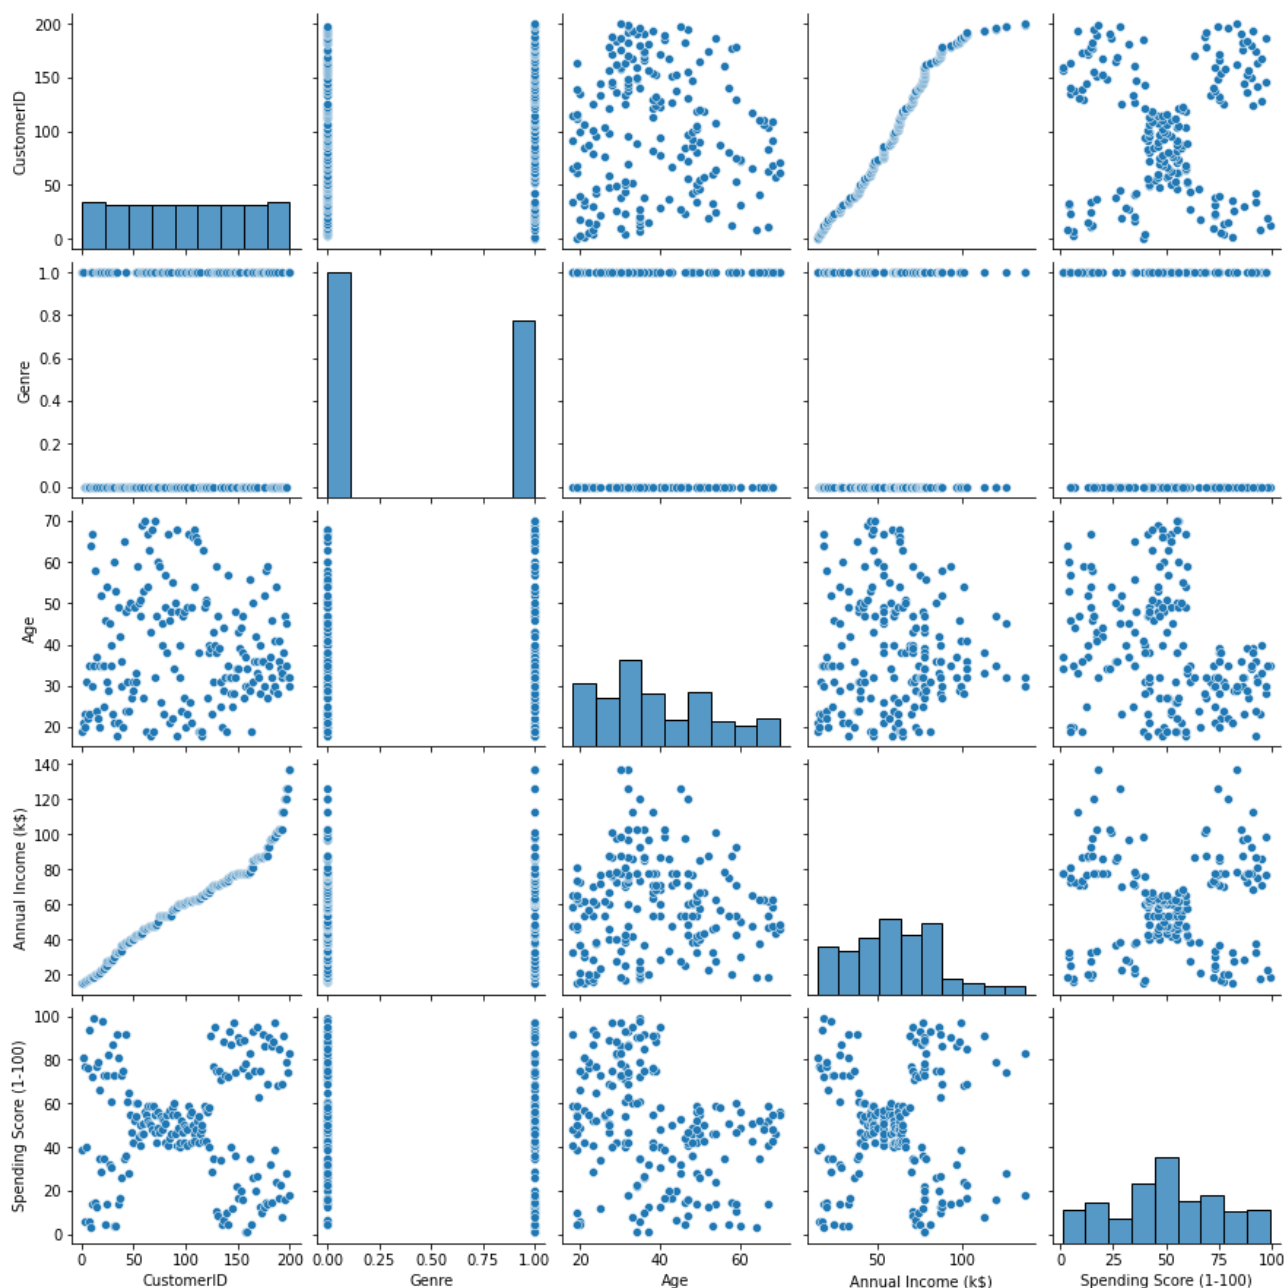
	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
114	115	0	18	65	48
111	112	0	19	63	54
115	116	0	19	65	50
2	3	0	20	16	6
39	40	0	20	37	75
...
102	103	1	67	62	59
108	109	1	68	63	43
57	58	1	69	44	46
60	61	1	70	46	56
70	71	1	70	49	55

200 rows × 5 columns

Step5 [Pair plot]

```
In [15]: 1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 sns.pairplot(data=customer_data)
```

Out[15]: <seaborn.axisgrid.PairGrid at 0x7f172d04f130>



Step6 [Build KMeans]

```
In [16]: 1 from sklearn.cluster import KMeans
```

```
In [17]: 1 customer_data.drop(['CustomerID'],axis=1, inplace=True)
```

```
In [18]: 1 KM = KMeans(n_clusters=5)
2 KM.fit(customer_data)
```

Out[18]: KMeans(n_clusters=5)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

1	KM.labels_
---	------------

[illegible]


```
1 print(KM.cluster_centers_)
```

```
[ [ 0.52777778 40.66666667 87.75      17.58333333]
  [ 0.4125     42.9375     55.0875     49.7125     ]
  [ 0.46153846 32.69230769 86.53846154 82.12820513]
  [ 0.39130435 45.2173913  26.30434783 20.91304348]
  [ 0.40909091 25.27272727 25.72727273 79.36363636] ]
```

Step7 [Scatter Plot]

```
1 sns.scatterplot(customer_data['Annual Income (k$)'], customer_data['Spending Sco
```

```
Out[21]: <AxesSubplot:xlabel='Annual Income (k$)', ylabel='Spending Score (1-100)'>
```



Step8 [Cluster Analysis]

```
1 kmeans2 = KMeans(n_clusters = 5, init='k-means++')
2 kmeans2.fit(customer_data)
3 pred = kmeans2.predict(customer_data)
```

```
1 frame = pd.DataFrame(customer_data)
2 frame['cluster'] = pred
```

```
In [24]: 1 frame.cluster.value_counts()
```

```
Out[24]: 1    79
          3    39
          2    37
          0    23
          4    22
          Name: cluster, dtype: int64
```

```
In [25]: 1 frame
```

```
Out[25]:
```

	Genre	Age	Annual Income (k\$)	Spending Score (1-100)	cluster
0	1	19	15	39	0
1	1	21	15	81	4
2	0	20	16	6	0
3	0	23	16	77	4
4	0	31	17	40	0
...
195	0	35	120	79	3
196	0	45	126	28	2
197	1	32	126	74	3
198	1	32	137	18	2
199	1	30	137	83	3

200 rows × 5 columns

```
In [26]: 1 C0 = customer_data[customer_data['cluster'] == 0]
          2 C1 = customer_data[customer_data['cluster'] == 1]
          3 C2 = customer_data[customer_data['cluster'] == 2]
          4 C3 = customer_data[customer_data['cluster'] == 3]
          5 C4 = customer_data[customer_data['cluster'] == 4]
```

```
In [27]: 1 import statistics as ss
          2 print('Average Age : ',C0['Age'].mean())
          3 print('Average Annual Income : ',C0['Annual Income (k$)'].mean())
          4 print('Deviation of the mean for annual Income : ',ss.stdev(C0['Annual Income (k$)']))
          5 print('No. of Customers ie shape : ',C0.shape)
          6 print('From those Customers We have',C0.Genre.value_counts()[1], 'males and',C0.Genre.value_counts()[0], 'females')
```

```
Average Age : 45.21739130434783
Average Annual Income : 26.304347826086957
Deviation of the mean for annual Income : 7.893811054517766
No. of Customers ie shape : (23, 5)
From those Customers We have 9 males and 14 females
```

```
In [28]: 1 print('Average Age : ',C1['Age'].mean())
          2 print('Average Annual Income : ',C1['Annual Income (k$)'].mean())
          3 print('Deviation of the mean for annual Income : ',ss.stdev(C1['Annual Income (k$)']))
          4 print('No. of Customers ie shape : ',C1.shape)
          5 print('From those Customers We have',C1.Genre.value_counts()[1], 'males and',C1.Genre.value_counts()[0], 'females')
```

```
Average Age : 43.12658227848101
Average Annual Income : 54.822784810126585
Deviation of the mean for annual Income : 8.576592314850398
No. of Customers ie shape : (79, 5)
From those Customers We have 33 males and 46 females
```

```
In [29]: 1 print('Average Age : ',C2['Age'].mean())
2 print('Average Annual Income : ',C2['Annual Income (k$)'].mean())
3 print('Deviation of the mean for annual Income : ',ss.stdev(C2['Annual Income (k$)'])
4 print('No. of Customers ie shape : ',C2.shape)
5 print('From those Customers We have',C2.Genre.value_counts()[1], 'males and',C2.Genre.value_counts()[0], 'females')
```

Average Age : 40.32432432432432
Average Annual Income : 87.43243243243244
Deviation of the mean for annual Income : 16.2729163891359
No. of Customers ie shape : (37, 5)
From those Customers We have 19 males and 18 females

```
In [30]: 1 print('Average Age : ',C3['Age'].mean())
2 print('Average Annual Income : ',C3['Annual Income (k$)'].mean())
3 print('Deviation of the mean for annual Income : ',ss.stdev(C3['Annual Income (k$)'])
4 print('No. of Customers ie shape : ',C3.shape)
5 print('From those Customers We have',C3.Genre.value_counts()[1], 'males and',C3.Genre.value_counts()[0], 'females')
```

Average Age : 32.69230769230769
Average Annual Income : 86.53846153846153
Deviation of the mean for annual Income : 16.312484972924967
No. of Customers ie shape : (39, 5)
From those Customers We have 18 males and 21 females

```
In [31]: 1 print('Average Age : ',C4['Age'].mean())
2 print('Average Annual Income : ',C4['Annual Income (k$)'].mean())
3 print('Deviation of the mean for annual Income : ',ss.stdev(C4['Annual Income (k$)'])
4 print('No. of Customers ie shape : ',C4.shape)
5 print('From those Customers We have',C4.Genre.value_counts()[1], 'males and',C4.Genre.value_counts()[0], 'females')
```


Average Age : 25.272727272727273
Average Annual Income : 25.727272727272727
Deviation of the mean for annual Income : 7.566730552584204
No. of Customers ie shape : (22, 5)
From those Customers We have 9 males and 13 females

Step9 [Find the best number of clusters]

```
In [32]: 1 SSE = []
2 for clust in range(1,20):
3     KM = KMeans(n_clusters= clust, init='k-means++')
4     KM = KM.fit(customer_data)
5     SSE.append(KM.inertia_)
```

```
In [33]: 1 plt.plot(np.arange(1,20), SSE, 'ro-')
2 plt.xlabel('Number of Clusters')
3 plt.ylabel('Inertia')
```

Out[33]: Text(0, 0.5, 'Inertia')



Step10 [Reduce Dimensions using PCA]

```
In [34]: 1 from sklearn.decomposition import PCA
```

```
In [35]: 1 pca = PCA(n_components=2)
2 _PCA = pca.fit_transform(customer_data)
3 PCA_Components = pd.DataFrame(_PCA)
```

```
In [36]: 1 PCA_Components
```

```
Out[36]:
```

	0	1
0	-31.534390	-33.349296
1	1.470393	-56.847684
2	-57.281081	-13.768454
3	-1.498970	-53.517338
4	-31.869456	-30.739696
...
195	58.010363	31.673893
196	19.158738	66.666049
197	58.097007	39.003394
198	20.086571	79.606665
199	71.977659	42.628194

200 rows × 2 columns


```
In [37]: 1 KM1 = KMeans(n_clusters=5)
          2 KM1.fit(PCA_Components)
          3 KM1.cluster_centers_
```

```
Out[37]: array([[ -44.30512124, -10.54356817],
                [ -10.68652483,  42.20770277],
                [  41.54479999,   2.33729767],
                [  -4.3966484 ,  -3.14441008],
                [   5.57443644, -46.63058222]])
```


```
In [38]: 1 KM1.labels_
```

[illegible]

Step11 [Scatter plot]

```
In [39]: 1 sns.scatterplot(PCA_Components[0], PCA_Components[1], hue=KM1.labels_)
```

```
Out[39]: <AxesSubplot:xlabel='0', ylabel='1'>
```



Step12 [MeanShift clustering]


```
In [40]: 1 from sklearn.cluster import MeanShift, AgglomerativeClustering
```

```
In [41]: 1 MS = MeanShift(bandwidth = 50)
          2 MS.fit(PCA_Components)
          3 MS.cluster_centers_
```

```
Out[41]: array([[ 0.3942943, -4.10083949]])
```

```
In [42]: 1 sns.scatterplot(PCA_Components[0], PCA_Components[1], hue=MS.labels_)
```

```
Out[42]: <AxesSubplot:xlabel='0', ylabel='1'>
```



Step13 [Predict hierarchical clusters using AgglomerativeClustering]

```
In [43]: 1 AC = AgglomerativeClustering(n_clusters = 5, linkage='ward',compute_full_tree=True)
2 AC.fit(customer_data)
```

```
Out[43]: AgglomerativeClustering(compute_full_tree=True, n_clusters=5)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [44]: 1 AC.labels_
```

```
Out[44]: array([4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3,
4, 3, 4, 3, 4, 0, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 0,
4, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 2, 1, 2, 1, 2, 1, 2,
0, 2, 1, 2, 1, 2, 1, 2, 1, 2, 0, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
1, 2])
```


```
In [45]: 1 customer_data['Cluster'] = AC.labels_
```

```
In [46]: 1 import scipy.cluster.hierarchy as sch
```

```
In [47]: 1 from scipy.cluster import hierarchy
```

In [48]:

```
1 Z = hierarchy.linkage(customer_data[:30], 'ward')
2 plt.figure(figsize=(10,5))
3 dn = hierarchy.dendrogram(Z)
```




Step14 [Visualize scatter plot with hue as agglomerativeclustering labels_]

In [49]:

```
1 sns.scatterplot(customer_data['Annual Income (k$)'], customer_data['Spending Sco
```

Out[49]: <AxesSubplot:xlabel='Annual Income (k\$)', ylabel='Spending Score (1-100)'



In [49]:

```
1
```