


▼ Step 2 : Importing Dataset

SURUTHI S
225229141

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
data = pd.read_csv("/content/fueldata.csv")
```

```
data.head()
```

	drivenKm	fuelAmount	
0	390.0	3600.0	
1	403.0	3705.0	
2	396.5	3471.0	
3	383.5	3250.5	
4	321.1	3263.7	

```
data.shape
```

```
(19, 2)
```

```
data.columns
```

```
Index(['drivenKm', 'fuelAmount'], dtype='object')
```

```
data.dtypes
```

```
drivenKm      float64
fuelAmount    float64
dtype: object
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19 entries, 0 to 18
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   drivenKm    19 non-null    float64
1   fuelAmount  19 non-null    float64
```

```
dtypes: float64(2)
memory usage: 432.0 bytes
```

▼ STEP 3: Pre Processing


```
data.isnull().sum()
```

```
drivenKm      0
fuelAmount    0
dtype: int64
```

▼ Step 4 : Vizualize Relationship

```
sns.relplot(data = data,x=data.drivenKm,y=data.fuelAmount)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f3d2795ffa0>
```



▼ STEP 5 Prepare X Matrix and Y vector

```
feature_list = data[['drivenKm']]
```

```
feature_list
```


	drivenKm	
0	390.00	
1	403.00	
2	396.50	
3	383.50	
4	321.10	
5	391.30	
6	386.10	
7	371.80	
8	404.30	
9	392.20	
10	386.43	
11	395.20	
12	381.00	
13	372.00	
14	397.00	
15	407.00	
16	372.40	
17	375.60	
18	399.00	

```
label =data[['fuelAmount']]
```

```
label
```

	fuelAmount 
0	3600.0
1	3705.0
2	3471.0
3	3250.5
4	3263.7
5	3445.2
6	3679.0
7	3744.5
8	3809.0
9	3905.0
10	3874.0
11	3910.0

```
data.describe()
```

	drivenKm	fuelAmount 
count	19.000000	19.000000
mean	385.548947	3710.684211
std	19.094297	281.892805
min	321.100000	3250.500000
25%	378.300000	3462.600000
50%	390.000000	3705.000000
75%	396.750000	3894.400000
max	407.000000	4235.900000

▼ Step 6 Examine X and Y

```
print(feature_list)
print("Type of X Matrix",type(feature_list))
print(label)
print("Type of Y Vector ",type(label))
```

```
drivenKm
```

```

0      390.00
1      403.00
2      396.50
3      383.50
4      321.10
5      391.30
6      386.10
7      371.80
8      404.30
9      392.20
10     386.43
11     395.20
12     381.00
13     372.00
14     397.00
15     407.00
16     372.40
17     375.60
18     399.00

```

Type of X Matrix <class 'pandas.core.frame.DataFrame'>

```

fuelAmount
0      3600.0
1      3705.0
2      3471.0
3      3250.5
4      3263.7
5      3445.2
6      3679.0
7      3744.5
8      3809.0
9      3905.0
10     3874.0
11     3910.0
12     4020.7
13     3622.0
14     3450.5
15     4179.0
16     3454.2
17     3883.8
18     4235.9

```

Type of Y Vector <class 'pandas.core.frame.DataFrame'>

➤ Step 7 Split dataset

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(feature_list, label,
                                                    test_size=0.20, random_state=42)
```

```
print(x_train, x_test, y_train, y_test)
```

```

drivenKm
8      404.30
16     372.40
3      383.50
13     372.00
15     407.00
17     375.60
2      396.50
9      392.20
18     399.00
4      321.10
12     381.00
7      371.80
10     386.43
14     397.00
6      386.10      drivenKm
0      390.0
5      391.3
11     395.2
1      403.0      fuelAmount
8      3809.0
16     3454.2
3      3250.5
13     3622.0
15     4179.0
17     3883.8
2      3471.0
9      3905.0
18     4235.9
4      3263.7
12     4020.7
7      3744.5
10     3874.0
14     3450.5
6      3679.0      fuelAmount
0      3600.0
5      3445.2
11     3910.0
1      3705.0

```

```
print(type(x_train))
```

```
<class 'pandas.core.frame.DataFrame'>
```

PART 1 :LR BASELINE MODEL

▼ STEP 8 BUILD MODEL

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

```
lin_reg.fit(x_train,y_train)
```

```
LinearRegression()
```

▼ STEP 9 PREDICT PRICE FOR 800 KM

```
lin_reg.predict([[800]])
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have
warnings.warn(
array([[6905.64571567]])
```

▼ STEP 10 PREDICT ON ENTIRE DATASET

```
from sklearn.metrics import mean_squared_error
yPred = lin_reg.predict(x_test)
```

```
yPred
```

```
array([[3775.81615646],
       [3785.74000628],
       [3815.51155575],
       [3875.05465468]])
```

▼ STEP 11 MSE

```
mse = mean_squared_error(y_test,yPred)
```

```
mse
```

```
46181.36710639155
```

```
lin_reg.coef_
```

```
array([[7.63373063]])
```

```
lin_reg.intercept_

array([798.6612099])

y_pred_data = lin_reg.predict(x_train)
y_pred_data

array([[3884.9785045 ],
       [3641.46249733],
       [3726.19690735],
       [3638.40900508],
       [3905.58957721],
       [3665.89043536],
       [3825.43540557],
       [3792.61036385],
       [3844.51973215],
       [3249.8521159 ],
       [3707.11258077],
       [3636.88225895],
       [3748.5637381 ],
       [3829.25227089],
       [3746.04460699]])
```

```
lin_reg.score(x_test,y_test)

-0.6180990161577022
```

PART 2 - LR WITH SCALING USING STANDARD SCALER (STANDARDIZATION)

▼ STEP 12 NORMALIZE USING STANDARD SCALER

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

norm_x_train = scaler.fit_transform(x_train)
norm_y_train = scaler.fit_transform(y_train)
norm_x_test = scaler.transform(x_test)
norm_y_test = scaler.transform(y_test)
```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:493: FutureWarning: The feature r


```
Feature names unseen at fit time:  
- drivenKm  
Feature names seen at fit time, yet now missing:  
- fuelAmount  
  
warnings.warn(message, FutureWarning)
```

▼ STEP 13 BUILD LR MODEL

```
norm_lreg = LinearRegression()  
  
norm_lreg.fit(norm_x_train,norm_y_train)  
  
LinearRegression()  
  
norm_yPred = norm_lreg.predict(norm_x_test)
```

▼ STEP 14 MSE

```
norm_mse = mean_squared_error(norm_y_test,norm_yPred)  
  
norm_mse  
  
32.25557286448923
```

▼ STEP 15 SCATTER PLOT

```
plt.plot(norm_y_test,norm_yPred,"go")
```

```
[<matplotlib.lines.Line2D at 0x7f3d23e89700>]
```



PART 3 - LR WITH SCALING USING MinMax SCALER (NORMALIZATION)

```
-5.850 |
```

▼ STEP 16 NORMALIZING USING MINMAX SCALER

```
from sklearn.preprocessing import MinMaxScaler
```

```
minmax = MinMaxScaler()
```

```
mm_norm_x_train = minmax.fit_transform(x_train)
```

```
mm_norm_y_train = minmax.fit_transform(y_train)
```

```
mm_norm_x_test = minmax.transform(x_test)
```

```
mm_norm_y_test = minmax.transform(y_test)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:493: FutureWarning: The feature r
```


```
Feature names unseen at fit time:
```

```
- drivenKm
```

```
Feature names seen at fit time, yet now missing:
```

```
- fuelAmount
```

```
warnings.warn(message, FutureWarning)
```



```
mm_norm_lreg = LinearRegression()
```

```
mm_norm_lreg.fit(mm_norm_x_train, mm_norm_y_train)
```

```
LinearRegression()
```

```
mm_norm_y_pred = mm_norm_lreg.predict(mm_norm_x_test)
```

```
mm_norm_y_pred
```

```
array([[ -1.93238929],
       [ -1.93151139],
       [ -1.92887767],
       [ -1.92361023]])
```

```
mm_norm_y_test
```

```
array([[0.3546783 ],
       [0.19758474],
       [0.66927136],
       [0.46123402]])
```

```
minmax_norm_mse = mean_squared_error(mm_norm_y_test,
                                       mm_norm_y_pred)
```

```
minmax_norm_mse
```

```
5.550397233153768
```

▼ prepare the model with input scaling

```
pipeline = Pipeline(steps=[('normalize', MinMaxScaler()), ('model', LinearRegression())])
```

```
fit pipeline
```

```
pipeline.fit(train_x, train_y)
```


```
make predictions
```

```
yhat = pipeline.predict(test_x)
```

```
y_test = y_test.to_numpy()
```


```
plt.xlabel("y_test")
plt.xlabel("y_pred")
plt.plot(y_test, yPred, "ro")
```

```
[<matplotlib.lines.Line2D at 0x7f3d23df3be0>]
```



```
plt.plot(norm_y_test,norm_yPred,"go")
```

```
[<matplotlib.lines.Line2D at 0x7f3d23dcf340>]
```



▼ STEP 17 KNN REGRESSOR

```
from sklearn.neighbors import KNeighborsRegressor
```

```
# creating Instance for the model
```

```
knn = KNeighborsRegressor(n_neighbors=5)
```

```
# Training / Fitting Data
```

```
knn.fit(x_train,y_train)
```

```
KNeighborsRegressor()
```

```
print(knn.predict([[800]]))
```

```
[[3829.08]]
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have  
warnings.warn(
```



```
knn_y_pred = knn.predict(x_test)
```

```
knn_mse = mean_squared_error(knn_y_pred,y_test)
```

```
knn_mse
```

21241.836200000045

▼ STEP 18 SGD REGRESSOR

```
from sklearn import linear_model
```

```
from sklearn.linear_model import SGDRegressor
```

```
from sklearn.pipeline import make_pipeline
```

```
max_iter = np.ceil(10**6/x_train.shape[0])
```

```
sgd = make_pipeline(StandardScaler(),linear_model.  
                    SGDRegressor(max_iter = max_iter,tol=1e-3))
```

```
print(type(x_train))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
x_train = x_train.to_numpy()
```

```
y_train = y_train.to_numpy()
```

```
sgd.fit(x_train,y_train)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning:
  y = column_or_1d(y, warn=True)
Pipeline(steps=[('standardscaler', StandardScaler()),
                 ('sgdregressor', SGDRegressor(max_iter=66667.0))])
```

```
sgd_y_pred = sgd.predict(x_test)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature r
  warnings.warn(
```

```
sgd_y_pred
```

```
array([3775.49866169, 3785.4202774 , 3815.18512453, 3874.71481879])
```

```
sgd_mse = mean_squared_error(y_test,sgd_y_pred)
```

```
sgd_mse
```

```
46085.64943360797
```

▼ STEP 19 SELECTING THE BEST MODEL

```
from tabulate import tabulate
```

```
data = [ ["MODELS","MSE VALUE"],  
         ["LINEAR REGRESSION",round(mse)],  
         ["STANDARD SCALER LR ",round(norm_mse)],  
         [" MINMAX  LR",round(minmax_norm_mse) ],  
         ["KNN",round(knn_mse)],  
         ["SGD",round(sgd_mse)]]
```

```
print(tabulate(data))
```

```
-----  
MODELS      MSE VALUE  
LINEAR REGRESSION  46181  
STANDARD SCALER LR  32  
MINMAX  LR      6  
KNN            21242  
SGD            46086  
-----
```

LINEAR REGRESSION MODEL AFTER NORMALIZING USING MINMAX SCALAR WOULD BE THE BEST MODEL AND HAS LOWEST MSE VALUE

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 8:43 AM

